Peng Qi

peng.qi@jd.com • pengqi@cs.stanford.edu • http://qipeng.me/

EDUCATION &
PROFESSIONAL
EXPERIENCES

Amazon AWS AI   Seattle, WA
Applied Scientist   Apr. 2022 – present

JD AI Research   Mountain View, CA
Senior Research Scientist   Jul. 2021 – Apr. 2022
Research Scientist   Oct. 2020 – Jun. 2021

Stanford University   Stanford, CA
Doctor of Philosophy, Computer Science   Sep. 2015 – Sep. 2020
- Research Assistant with Stanford Natural Language Processing Group (Advisor: Chris Manning)

Master of Science, Computer Science   Sep. 2013 – Jun. 2016
- Research Assistant with Stanford Artificial Intelligence Laboratory (Advisor: Andrew Ng)

Master of Science, Department of Statistics   Apr. 2016 – Apr. 2017

Tsinghua University   Beijing, China
Research Assistant, State Key Laboratory of Intelligent   Jul. 2011 – Jun. 2013
         Technology & Systems (Advisor: Xiaolin Hu)
Bachelor of Engineering (magna cum laude), School of Software   Aug. 2008 – Jul. 2012

PUBLICATIONS
(* = Equal
Contribution)

[1] Chao Shang, Guangtao Wang, **Peng Qi**, and Jing Huang. Improving Time Sensitivity for Question Answering over Temporal Knowledge Graphs. In *Association of Computational Linguistics (ACL)*, 2022.

[2] **Peng Qi***, Haejun Lee*, Oghenetegiri "TG" Sido*, and Christopher D. Manning. Answering Open-Domain Questions of Varying Reasoning Steps from Text. In *Empirical Methods for Natural Language Processing (EMNLP)*, 2021.

[3] Chao Shang, **Peng Qi**, Guangtao Wang, Jing Huang, Youzheng Wu, Bowen Zhou. Open Temporal Relation Extraction for Question Answering. In *3rd Conference on Automated Knowledge Base Construction (AKBC)*, 2021.

[4] Kevin Huang, **Peng Qi**, Guangtao Wang, Tengyu Ma, and Jing Huang. Entity and Evidence Guided Document-Level Relation Extraction. In *6th Workshop on Representation Learning for NLP (RepL4NLP) at ACL 2021*, 2021.

[5] Xiaochen Hou, Jing Huang, Guangtao Wang, **Peng Qi**, Xiaodong He, and Bowen Zhou. Selective Attention Based Graph Convolutional Networks for Aspect-Level Sentiment Classification. In *TextGraphs-15 at NAACL 2021*, 2021.

[6] Yuhao Zhang, Yuhui Zhang, **Peng Qi**, Christopher D. Manning, and Curtis P. Langlotz. Biomedical and Clinical English Model Packages for the Stanza Python NLP Library. *Journal of the American Medical Informatics Association (JAMIA)*, 2021.

[7] Xiaochen Hou, **Peng Qi**, Guangtao Wang, Rex Ying, Jing Huang, Xiaodong

He, and Bowen Zhou. Graph Ensemble Learning over Multiple Dependency Trees for Aspect-level Sentiment Classification. In *2021 Annual Meeting of the North American Chapter of the Association for Computational Linguistics (NAACL)*, 2021.

[8] Devendra Singh Sachan, Yuhao Zhang, **Peng Qi**, and William L. Hamilton. Do Syntax Trees Help Pretrained Transformers Extract Information? In *The 16th Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, 2021.

[9] Ashwin Paranjape*, Abigail See*, Kathleen Knealy, Haojun Li, Amelia Hardy, **Peng Qi**, Kaushik Ram Sadagopan, Nguyet Minh Phu, Dilara Soylu, and Christopher D. Manning. Neural Generation Meets Real People: Towards Emotionally Engaging Mixed-Initiative Conversations. *The Alexa Prize Proceedings*, 2020.

[10] **Peng Qi**. Explainable and Efficient Knowledge Acquisition from Text. (Ph.D. Thesis) *Stanford University*, 2020.

[11] **Peng Qi**, Yuhao Zhang, and Christopher D. Manning. Stay Hungry, Stay Focused: Generating Informative and Specific Questions in Information-Seeking Conversations. In *Findings of ACL: EMNLP 2020*, 2020.

[12] **Peng Qi***, Yuhao Zhang*, Yuhui Zhang, Jason Bolton, and Christopher D. Manning. Stanza: A Python Natural Language Processing Toolkit for Many Human Languages. In *Association of Computational Linguistics (ACL), System Demonstrations*, 2020

[13] **Peng Qi**, Xiaowen Lin*, Leo Mehr*, Zijian Wang*, Christopher D. Manning. Answering Complex Open-domain Questions Through Iterative Query Generation. In *2019 Conference on Empirical Methods in Natural Language Processing and 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP), 2019.*

[14] Zhilin Yang*, **Peng Qi***, Saizheng Zhang*, Yoshua Bengio, William W. Cohen, Ruslan Salakutdinov, and Christopher D. Manning. HotpotQA: A Dataset for Diverse, Explainable Multi-hop Question Answering. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2018.

[15] Yuhao Zhang*, **Peng Qi***, and Christopher D. Manning. Graph Convolution over Pruned Dependency Trees Improves Relation Extraction. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2018.

[16] **Peng Qi***, Timothy Dozat*, Yuhao Zhang*, and Christopher D. Manning. Universal Dependency Parsing from Scratch. In *CoNLL 2018 Shared Task: Multilingual Parsing from Raw Text to Universal Dependencies*, 2018.

[17] Urvashi Khandelwal, He He, **Peng Qi**, and Dan Jurafsky. Sharp Nearby, Fuzzy Far Away: How Neural Language Models Use Context. In 56^{th} *Annual Conference of Association of Computational Linguistics (ACL)*, 2018.

[18] **Peng Qi** and Christopher D. Manning. Arc-swift: A Novel Transition System for Dependency Parsing. In 55^{th} *Annual Conference of Association of Computational Linguistics (ACL)*, 2017.

[19] Timothy Dozat, **Peng Qi**, and Christopher D. Manning. Stanford's Graph-based Neural Dependency Parser at the CoNLL 2017 Shared Task. *CoNLL 2017 Shared*

Task: Multilingual Parsing from Raw Text to Universal Dependencies. *First place*

[20] Arun Chaganty*, Ashwin Paranjape*, Jason Bolton*, Matthew Lamm*, Jinhao Lei*, Abigail See*, Kevin Clark, Yuhao Zhang, **Peng Qi**, and Christopher D. Manning. Stanford at TAC KBP 2017: Building a Trilingual Relational Knowledge Graph. In *Text Analysis Conference (TAC)*, 2017.

[21] Yuhao Zhang*, Arun Chaganty*, Ashwin Paranjape*, Danqi Chen*, Jason Bolton*, **Peng Qi**, and Christopher D. Manning. Stanford at TAC KBP 2016: Sealing Pipeline Leaks and Understanding Chinese. In *Proceedings or the Text Analysis Conference (TAC) - Knowledge Base Population*, 2016.

[22] Andrew L. Maas, **Peng Qi**, Ziang Xie, Awni Y. Hannun, Daniel Jurafsky, and Andrew Y. Ng. Building DNN Acoustic Models for Large Vocabulary Speech Recognition. *Computer Speech & Language*, 2016.

[23] **Peng Qi** and Xiaolin Hu. Learning nonlinear statistical regularities in natural images by modeling the outer product of image intensities. *Neural computation*, 26(4):693–711, 2014.

[24] Xiaolin Hu, Jianwei Zhang, **Peng Qi**, and Bo Zhang. Modeling response properties of V2 neurons using a hierarchical k-means model. *Neurocomputing*, 134:198–205, 2014.

[25] **Peng Qi**, Shuochen Su, and Xiaolin Hu. Modeling outer products of features for image classification. In *Advanced Computational Intelligence (ICACI)*, 2013.

[26] Xiaolin Hu, **Peng Qi**, and Bo Zhang. Hierarchical k-means algorithm for modeling visual area V2 neurons. In *Neural Information Processing (ICONIP)*, pages 373–381, 2012. *Best Paper Award*

Preprints & In submission

[27] **Peng Qi**, Jing Huang, Youzheng Wu, Xiaodong He, and Bowen Zhou. Conversational AI Systems for Social Good: Opportunities and Challenges. *arXiv preprint arXiv:2105.06457*, 2021.

[28] Bo Li, **Peng Qi**, Bo Liu, Shuai Di, Jingen Liu, Jiquan Pei, Jinfeng Yi, and Bowen Zhou. Trustworthy AI: From Principles to Practices. *arXiv preprint arXiv:2110.01167*, 2021.

[29] **Peng Qi***, Guangtao Wang*, and Jing Huang. SpanDrop: Simple and Effective Counterfactual Learning for Long Sequences. Submitted to *International Conference on Machine Learning*, 2022.

HONORS   Inaugural Yufan Award (Rising Stars Category) at World Artificial Intelligence Conference 2020 (20 top young Chinese AI scientists worldwide), Facebook ParlAI Research Award, China's National Scholarship (top 3% university-wide at Tsinghua), Freshman Scholarship (provincial top 10 in college entrance exam), and other merit-based awards from undergrad

EXPERIENCE   Facebook AI Research (New York)   Jun. 2017 – Sep. 2017
Learning to Teach through Communication (Mentors: Jason Weston, Douwe Kiela, Kyunghyun Cho)

- Studied emergent teaching behavior of machine learning agents in a constrained communication setting
- Proposed a reinforcement learning-based method for agents to learn to teach
- Implemented the proposed method in ParlAI with positive results on image classification tasks

Contributing Open Source Projects

Stanza ⬈ (documentation ⬈)
- PyTorch implementation of Stanford's full system in the 2018 CoNLL Shared Task on Universal Dependency Parsing from Raw Text
- Includes neural network models for tokenization, part-of-speech tagging, lemmatization, dependency parsing, and NER in 60+ languages with pretrained models available and intuitive Python interface
- Pipelines for syntactic analysis and NER for English-language biomedical text
- Main contributor and maintainer

HuggingFace Transformers ⬈
- Bugfixes in various standard example scripts for running BERT for question answering

Universal Dependencies ⬈
- Maintainer of the Chinese-GSD treebank and its corresponding version in simplified Chinese
- Converted the GSD treebank to simplified Chinese, fixed annotations and segmentation when necessary, engaged in community discussions about annotation standards

ParlAI ⬈
- Implemented various features to enable teacher training in ParlAI
- Bugfixes and stability improvements

Caffe ⬈
- Implemented a framework for generic solvers
- Implemented Nesterov's Accelerated Gradient solver and AdaGrad solver
- Contributed neuron layers (Leaky ReLU, Mean-variance normalization)
- Compatibility issues / bug fixes for Mac OS

Kaldi ⬈
- Contributed the first training recipe for the Fisher/Switchboard mixed speech corpus (the largest speech corpus in use in academia)
- A bug fix for Kaldi's speaker identification for better speaker heldout training

ConvolutionalRBM.m ⬈ (Owner)
- An implementation of Honglak Lee *et al.*'s convolutional restricted Boltzmann machine model in Matlab, MEX (C++/CUDA)

INVITED TALKS
- Seminar speaker. Explainable and Efficient Knowledge Acquisition from Text. University of Southern California. (Spring 2020)
- Seminar speaker. Explainable and Efficient Knowledge Acquisition from Text. University of Arizona. (Spring 2020)
- Seminar speaker. Answering Complex Open-domain Questions in the Wild. University of Southern California Information Sciences Institute (USC ISI). (Fall 2019)
- Seminar speaker. Answering Complex Open-domain Questions in the Wild. Stanford OVAL Seminar. (Fall 2019)
- Seminar speaker. Answering Complex Open-domain Questions in the Wild. University of California at Los Angeles. (Fall 2019)
- Invited speaker. Open-domain Text Understanding at Scale with Multi-step Reasoning. Facebook AI Research. (Summer 2019)

- Invited speaker. Open-domain Text Understanding at Scale with Multi-step Reasoning. Cresta, Inc. (Summer 2019)
- Invited speaker. Open-domain Text Understanding at Scale with Multi-step Reasoning. Samsung Research America. (Summer 2019)
- Invited speaker. Open-domain Text Understanding at Scale with Multi-step Reasoning. Salesforce Research. (Spring 2019)
- Seminar speaker. Answering Complex Open-domain Questions Through Iterative Query Generation. University of Washington. (Summer 2019)
- Speaker. A Gentle Introduction to Reinforcement Learning. Stanford NLP Reading Group. (Spring 2019)
- Invited speaker. A tutorial on Deep Learning for Natural Language Processing. LingCon Hackathon training day ↗ (Fall 2017)

Teaching &
Mentoring
- TA, CS 224D Deep Learning for Natural Language Processing (Spring 2015)
- TA, CS 124 From Languages to Information (Winter 2015)
- TA, CS 145 Introduction to Databases (Summer 2014)
- TA, CS 224S Spoken Language Processing (Spring 2014)
- Project Mentor, CS 224n (Winter 2019, Best Custom Project Report Prize; Winter 2018, Best Custom Project Report Prize; Winter 2017)
- Tutor, CS 145 Introduction to Databases, CS 107 Computer Organization and Systems, CS 245 Database Systems Principles

Service

Session Chair: NAACL 2021

Area Chair: NAACL 2021

Publicity Chair: NAACL 2021

Program Committee:
2021 ACL Rolling Review, AKBC, MRQA, AAAI
2020 EMNLP (Outstanding Reviewer)
2019 EMNLP-IJCNLP (Outstanding Reviewer), MRQA, ACL, NAACL-HLT
2018 EMNLP (Best Reviewer Award), ACL, CoNLL Shared Task, UDW
2017 ACL
2013 ICACI
Journal: IEEE TNNLS